

Lab #3

Purpose: In this lab, we implement the **logistic regression and KNN algorithms**.

1. Load Data Default.csv to object **Default**.

```
rm(list = ls())
```

```
Default=read.csv(file.choose(),header=T)
```

```
attach(Default)
```

attach allows you to reference database variables without using the "\$" extractor

```
names(Default)
```

```
summary(Default)
```

Here header is true because the CSV file has column headings in it

The `attach` function in R is a powerful tool in base R that allows you to, well, attach a specific database to a search path. In other words, `attach` allows you to reference database variables without using the "\$" extractor. For example, without using the `attach` function (using the Iris data set from R):

```
> head(iris)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1 5.1         3.5         1.4         0.2      setosa
2 4.9         3.0         1.4         0.2      setosa
3 4.7         3.2         1.3         0.2      setosa
4 4.6         3.1         1.5         0.2      setosa
5 5.0         3.6         1.4         0.2      setosa
6 5.4         3.9         1.7         0.4      setosa
```

```
> #calculating the mean
> mean(iris$Sepal.Length)
[1] 5.843333
```

Now, with the `attach` function:

```
> attach(iris)
> mean(Sepal.Length)
[1] 5.843333
```

2. Logistic Regression

The original data set contains **10000** observations. We first fit model **glm.fit1** which includes all three independent variables: student, balance and income, with all the observations. We can use the **glm()** function.

```
glm.fit1=glm(default~balance+income+student,family="binomial",data=Default)
```

We can use the function **summary()** to obtain the coefficients, their associated p-values, model AIC and residual deviance. The residual deviance is 1571.5 and the AIC is 1579.5. The function **exp()** calculates **odds ratio of the coefficients**.

```
summary(glm.fit1)
```

```
coef(glm.fit1)
```

```
exp(coef(glm.fit1))
```

We next evaluate the prediction accuracy of this model. We use the **sample()** function to split the original data set into one training set and one test set. We randomly select **80%** of the observations for training and the remaining in the test set (**Default.test**). And save the true value of the testing set in the vector **test.truevalue**.

```
set.seed(1)
```

```
train=sample(nrow(Default),nrow(Default)*0.8)
```

```
Default.test=Default[-train, ]
```

```
test.truevalue=default[-train]
```

```
train.truevalue=default[train]
```

We fit the data with all three independent variables into a logistic regression model **glm.fit2**. The **subset** option in **glm()** is to fit a regression using only the observations corresponding to the subset.

```
glm.fit2=glm(default~balance+student+income,data=Default,subset=train,family =binomial)
```

```
summary(glm.fit2)
```

```
exp(coef(glm.fit2))
```

We evaluate the performance of the model using the test set (`Default.test`). We use the function `predict()` to calculate the predicted probabilities of the default in the test set and store them to `glm.pred2`. The `type="response"` option tells R to output probabilities not the logit.

```
glm.probs2=predict(glm.fit2,Default.test, type="response")
```

We convert the predicted probabilities into a binary class label, `Yes` or `No`. The following commands create a vector of class predictions based on whether the predicted probability is greater than or less than 0.5.

```
glm.pred2=rep("No",2000)
glm.pred2[glm.probs2>.5]="Yes"
```

observation
if probability > 0.5 ⇒ Yes

The `table()` function produce a confusion matrix to determine how many observations were correctly classified. We can find that `glm.fit2` correctly predicts 97.25% of default status and misclassify 2.75% in the test set.

```
table(glm.pred2,test.truevalue)
mean(glm.pred2==test.truevalue)
```

the order does matter when the implication is meaningful

```
mean(glm.pred2!=test.truevalue)
```

`glm.fit2` correctly predicts 97.5% of default status.

$$\frac{1930 + 20}{2000} = 97.5\%$$

We next evaluate the 5-fold cross-validation prediction of the model.

```
k=5
```

```
folds=sample(1:k,nrow(Default),replace=TRUE)
```

We create a vector to store the accuracy result for each fold. We set the initial values for this vector as zero.

```
accuracy=rep(0,k)
```

```
for(i in 1:k)
{
  glm.fit3=glm(default~balance+student+income,family="binomial",data=Default[folds!=i,])
  Default.test=Default[folds==i, ]
```

```
  glm.probs3 =predict(glm.fit3,Default.test, type="response")
  glm.pred3=rep("No",nrow(Default[folds==i,]))
```

```
  glm.pred3[glm.probs3>.5]="yes"
```

```
  test.truevalue=default[folds==i]
  accuracy[i]=mean(glm.pred3==test.truevalue)
```

```
}
```

Then we calculate the average cross-validation accuracy, which is 96.3%.

```
mean(accuracy)
```

3. KNN

We perform KNN using the `knn()` function, which is part of the `class` package. We first load the `class` package.

```
library(class)
```

We first normalize the two quantitative input variables `balance` and `income` so that they would be on a comparable scale. The `scale()` function standardizes the quantitative variables.

```
standardized.balance=scale(balance)
```

```
standardized.income=scale(income)
```

Then we use the function `bind()` to combine the two standardized variables and the qualitative input variable together.

```
Input.standard=cbind(standardized.balance,standardized.income,student)
```

```
accuracy=matrix(0,10,5)
```

To use KNN, we need to determine K. We can use 5-fold cross-validation to select the best K from [1,10]. Thus, we first create a matrix to store the accuracy results for five folds and ten different K values. We set the initial values for this matrix as zero.

```
set.seed(2)
```

```
folds=sample(1:5,nrow(Input.standard),replace=TRUE)
```

```
for (j in 1:10)
{
  for(i in 1:5)
  {
    train.standard=Input.standard[folds!=i,]
    test.standard=Input.standard[folds==i,]
    train.truevalue=default[folds!=i]
    test.truevalue=default[folds==i]
    knn.pred=knn(train.standard,test.standard[train.truevalue,k=j])
    table(knn.pred,test.truevalue)
    accuracy[j,i]=mean(knn.pred==test.truevalue)
  }
}
```

`train.standard` is the input matrix of the training set and `test.standard` is the input matrix of the testing set. `train.truevalue` is the original value of default in the training set and `test.truevalue` is the true value of default in the testing set. For each observation in the testing set, the `knn` function can calculate its distance with each observation in the training set based on `train.standard` and `test.standard`. Given `k=j`, it selects the `k` most nearest neighbors, and use their default values (based on `train.truevalue`) to predict the default values in the testing set. In the `knn` function, we specify the number of neighbors in the option "`k=j`". The output of this loop is the prediction accuracy

Then we calculate the average cross-validation accuracy for each K. The function `which.max()` tells us that when `K=9`, it leads to the highest accuracy as 97.17% .Thus, K should be selected as 9.

```
cv.accuracy=apply(accuracy,1,mean)
```

```
which.max(cv.accuracy)
```

This concludes lab #3.